

Automated Pneumonia Detection from Chest X-Ray Images Using Machine Learning



Akash Kumar Yadav
MS in Computer Science
Fordham University

Problem Statement

The Global Health Challenge

- ▶  **Problem:** Pneumonia causes ~2.5 million deaths annually.
- ▶  **Impact:** Responsible for 15% of child mortality under age 5 (WHO).
- ▶  **Challenge:** Manual X-ray interpretation is time-consuming and error-prone.
- ▶  **Gap:** Expert radiologists are often unavailable in resource-limited settings.

Project Goal:

Develop an automated screening system achieving **>95% sensitivity**.

Related Work & Contributions

Prior Work

- ▶ **CheXNet (DenseNet)**: Achieved radiologist-level accuracy.
- ▶ **Transfer Learning**: Proven effective for medical imaging tasks.
- ▶ **Limitation**: Limited work exists on interpretability & robustness validation.

Our Contributions

- ✓ **Comparison**: Evaluated 3 architectures (Custom CNN, VGG16, ResNet50).
- ✓ **Performance**: Achieved **99.74% recall** (missing only 1 case per 390).
- ✓ **Interpretability**: Validated decision-making using Grad-CAM.
- ✓ **Robustness**: Independent validation confirming stability.

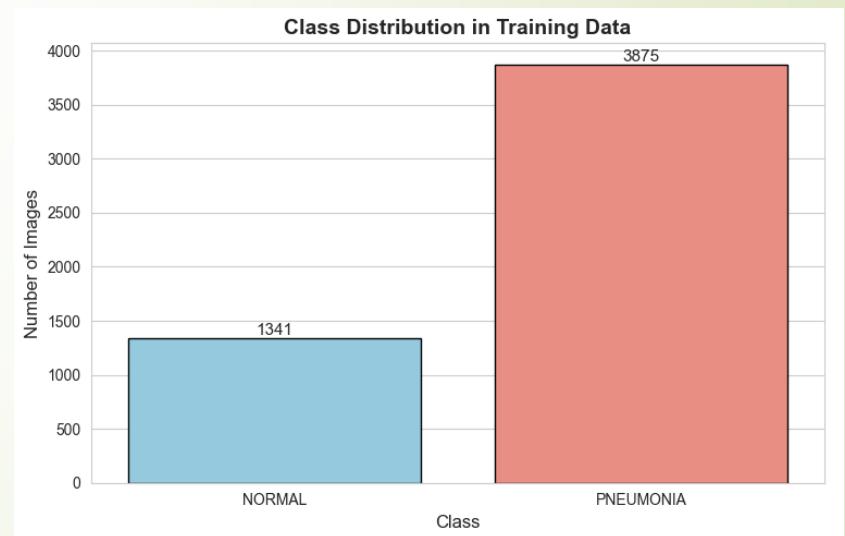
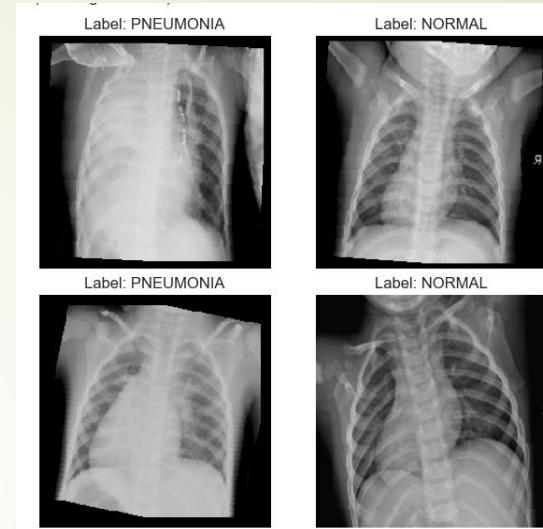
Dataset Overview

Chest X-Ray Dataset Breakdown

Total Images: 5,840 (Pediatric patients, 1-5 years)

- └ Training: 5,216 (89%)
 - └ Normal: 1,341 (25.7%)
 - └ Pneumonia: 3,875 (74.3%)
- └ Testing: 624 (11%)
 - └ Normal: 234 (38%)
 - └ Pneumonia: 390 (62%)

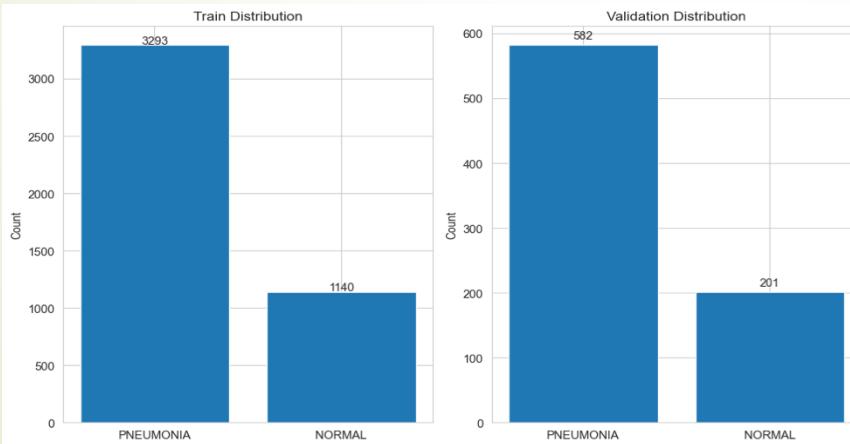
⚠ Class Imbalance: 3:1 ratio (Pneumonia : Normal)



Methodology & Data Augmentation

Training Strategy

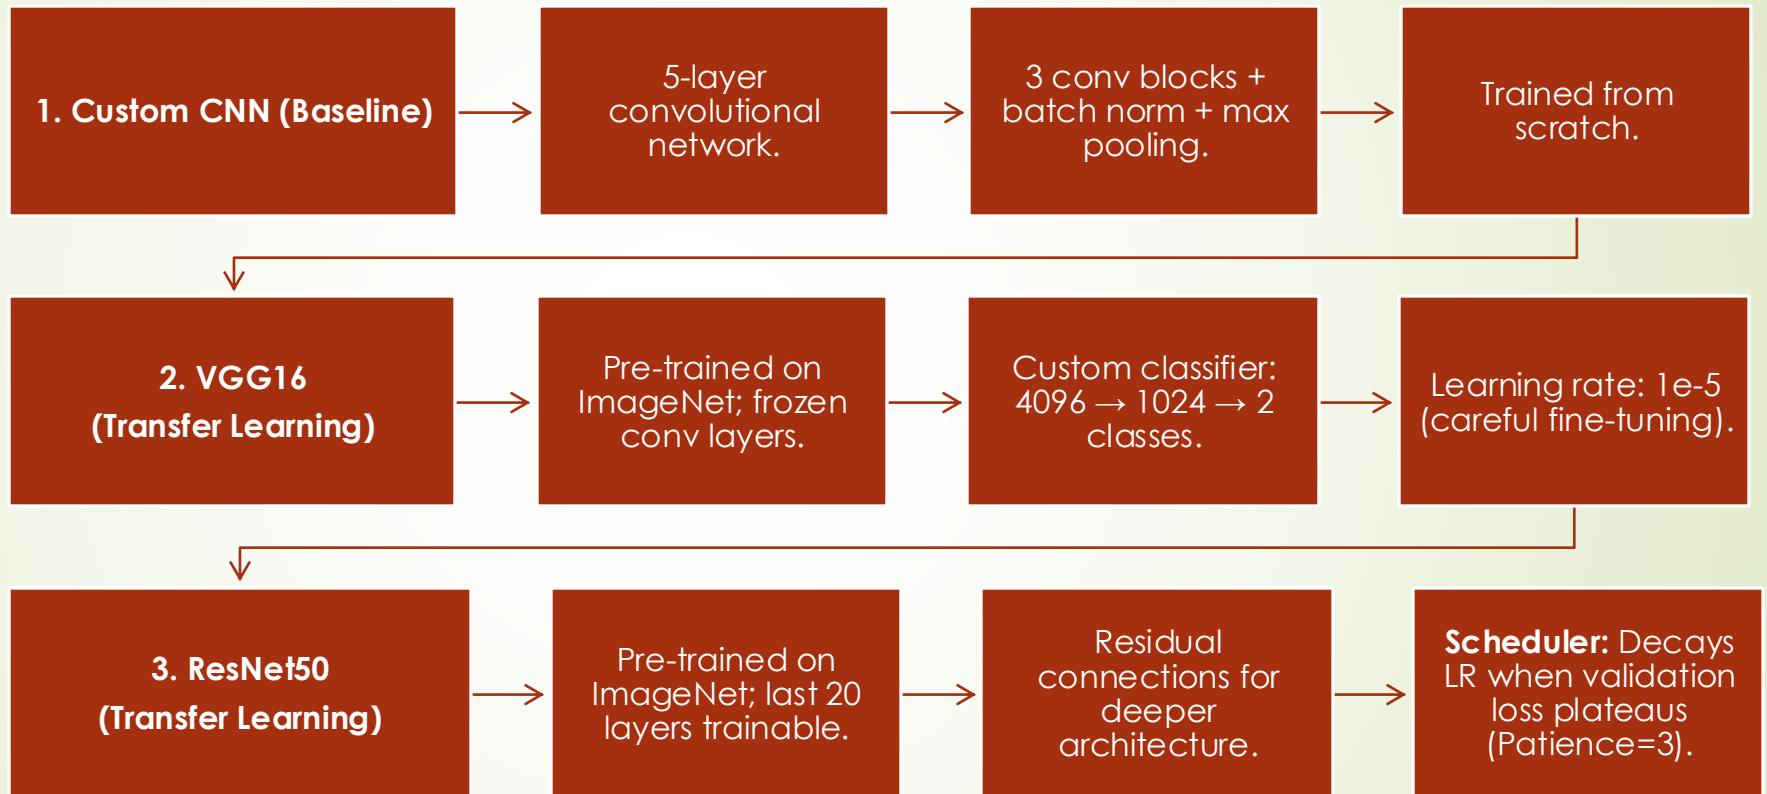
- **Splits:** 85% Train / 15% Validation (main); 80/20 (independent).
- **Loss Function:** Weighted Cross-Entropy (Weights: Normal ~1.94, Pneumonia ~0.67).
- **Optimizer:** Adam with ReduceLROnPlateau scheduler.



Data Augmentation Techniques

Technique	Parameters	Purpose
Rotation	$\pm 15^\circ$	Fix orientation issues
Horizontal Flip	$p=0.5$	Mirror views
Color Jitter	± 0.1 bright/cont rast	Lighting variation
Normalization	ImageNet mean/std	Standardization
Random Scaling	$0.9x - 1.1x$	Size invariance

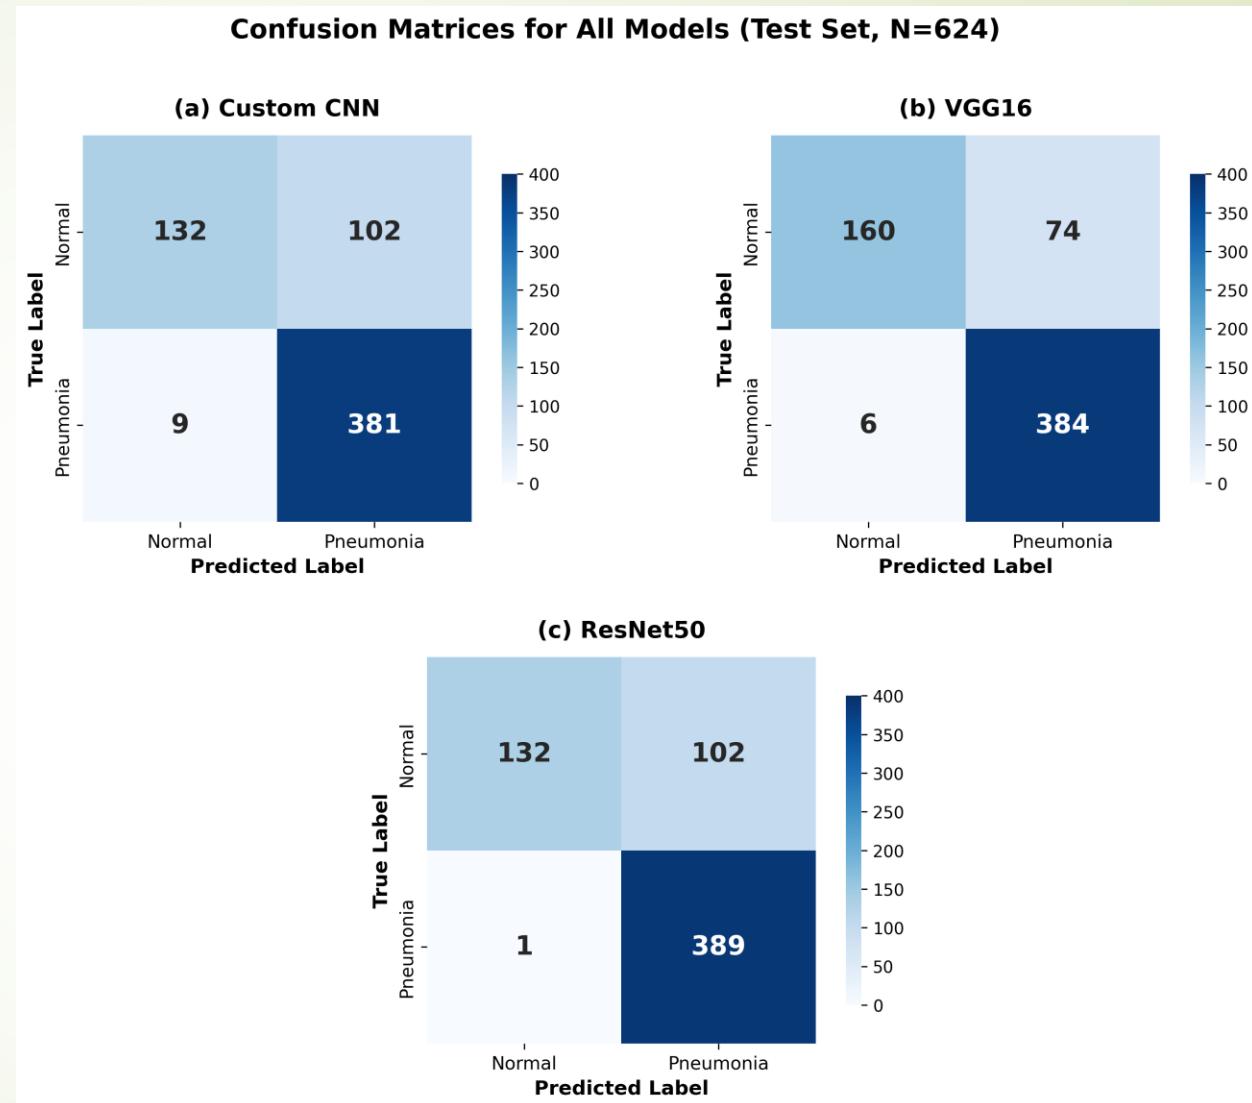
Model Architectures



Results - Confusion Matrix

► Key Finding:

ResNet50 missed only 1 pneumonia case out of 390!



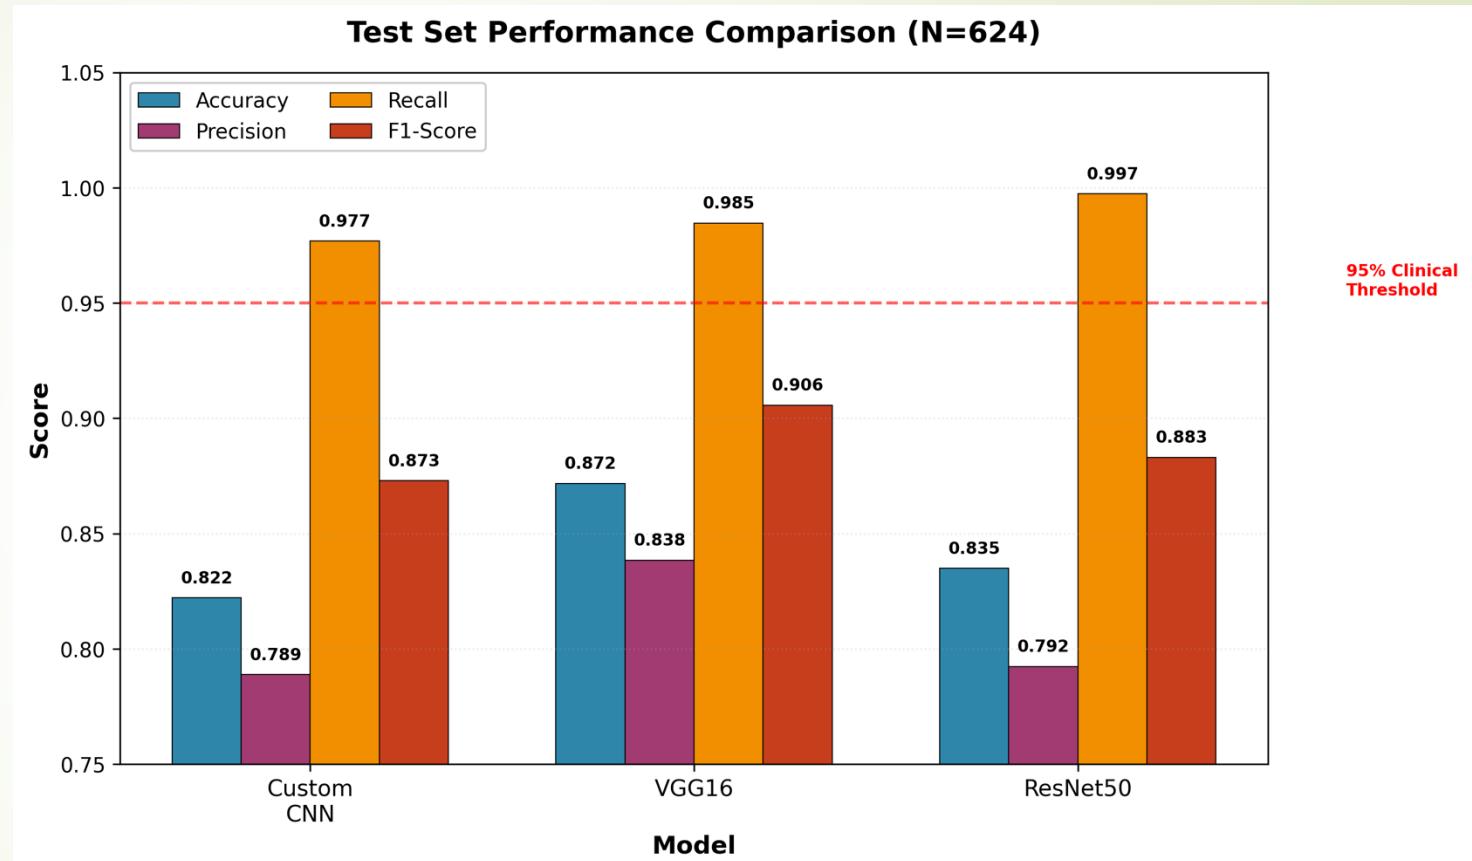
Results - Performance Metrics

► Key Finding

✓ All models exceed 95% clinical threshold.

ResNet50: Highest recall (99.74%).

VGG16: Best balanced F1-score (90.57%).

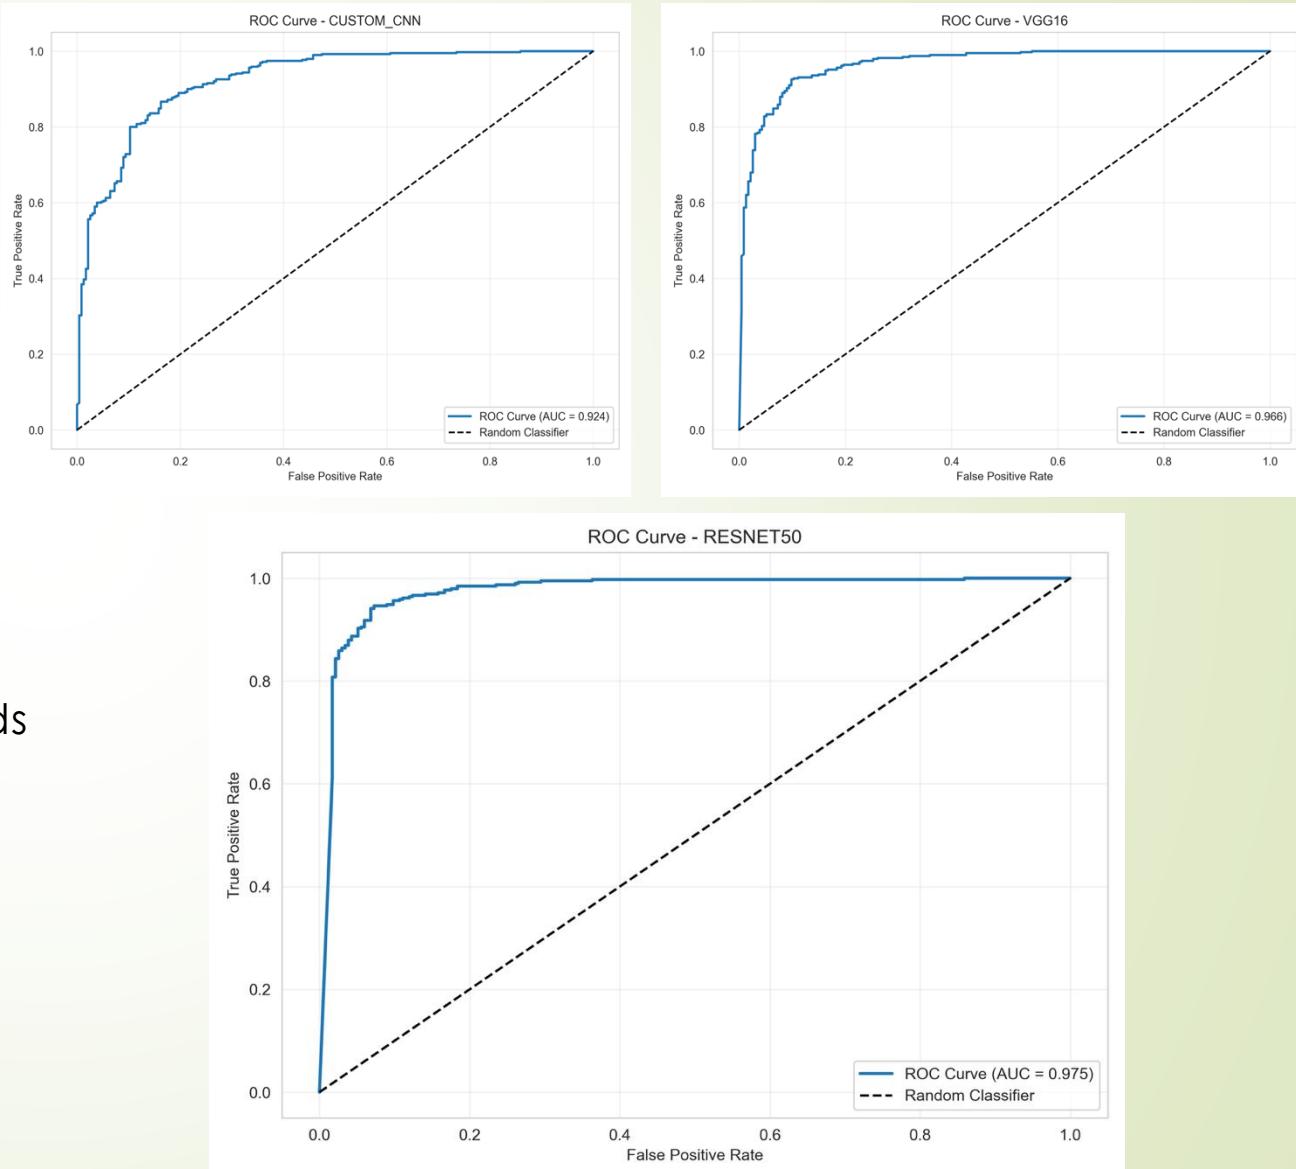


ROC Curves & AUC Analysis

ROC-AUC Scores (Higher is Better)

- ▶ **Custom CNN:** 92.4%
- ▶ **VGG16:** 96.60%
- ▶ **ResNet50:** 97.45% ★

📊 **Key Insight:** ResNet50 shows excellent discrimination across all decision thresholds (AUC 97.45%).

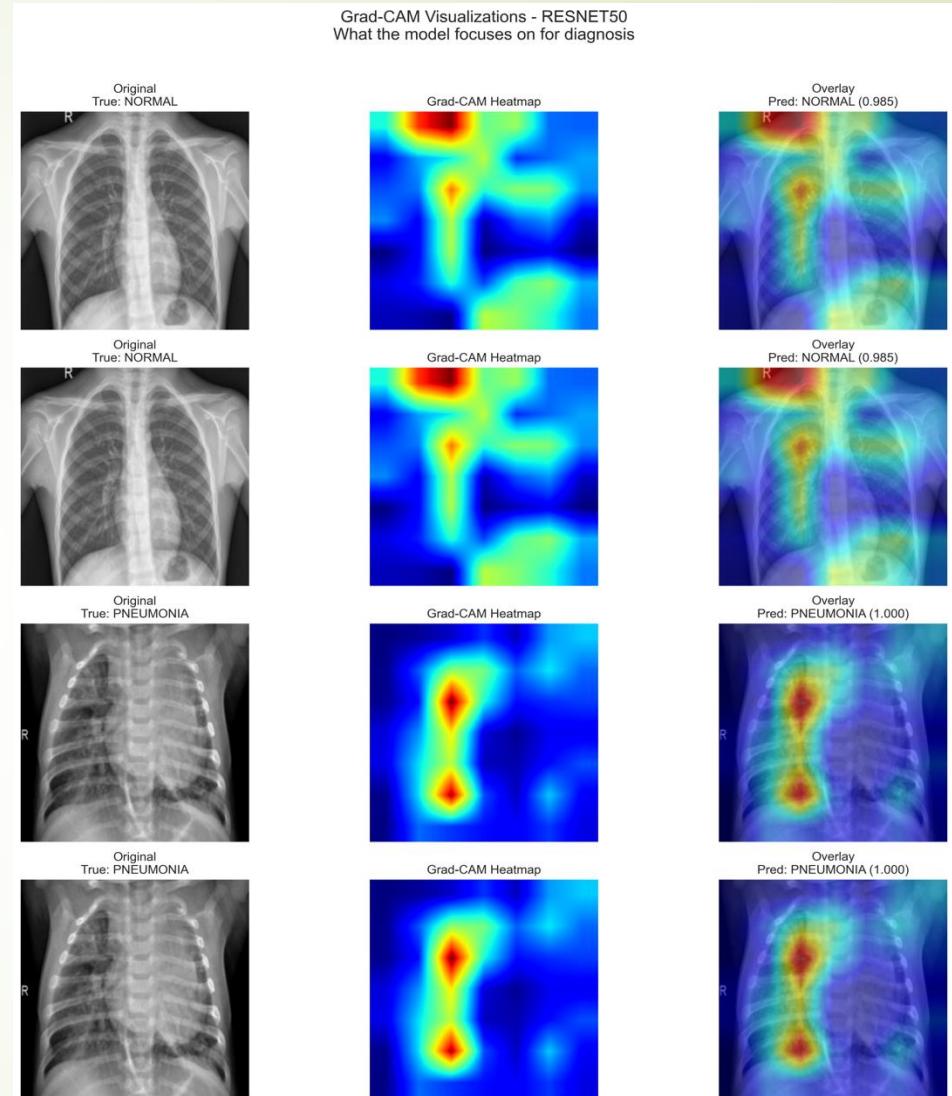


Interpretability – Grad-CAM

Grad-CAM Findings

- ✓ Models focus on **lung regions** (clinically relevant).
- ✓ Strong activation on **infiltrates** in pneumonia cases.
- ✓ Minimal activation in normal cases.
- ✓ **No learning** of image artifacts or borders.'

Result: High clinical trust via explainable decisions.



Clinical Significance & Deployment

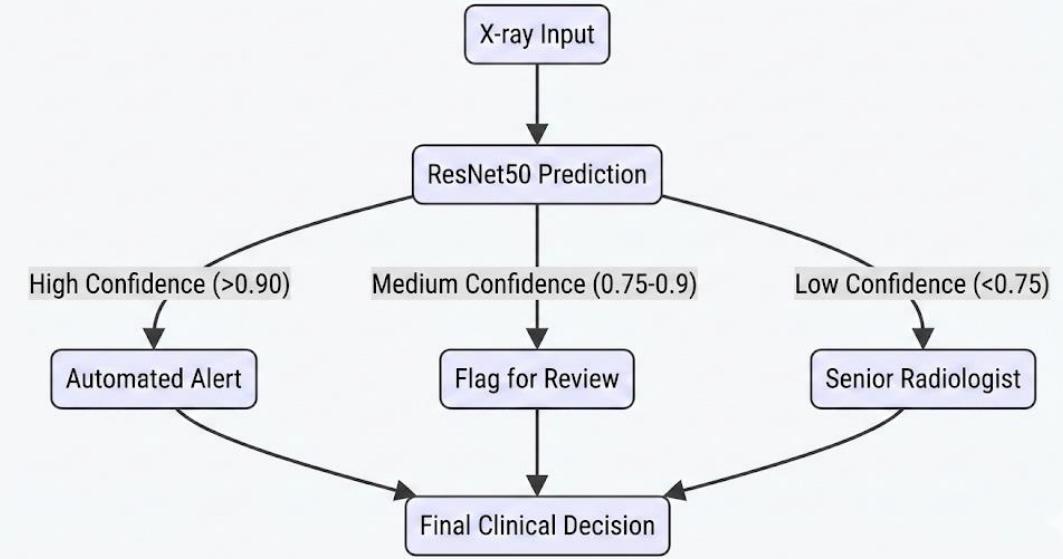
Clinical Impact

🎯 **0.26% miss rate** (Only 1 case missed per 390 patients).

🏥 Exceeds clinical sensitivity requirements.

🔄 Higher False Positive rate (40%) is acceptable for screening safety.

AI-Assisted X-ray Triage Process Flowchart



Conclusion & Future Work

Key Achievements

- ✓ **Best-in-class Performance:** 99.74% recall with ResNet50.
- ✓ **Validated Robustness:** 98.70% ($\pm 0.28\%$) via K-fold cross-validation.
- ✓ **Clinical Interpretability:** Grad-CAM validates clinically relevant features.

Future Directions

- 🌐 External validation on diverse datasets.
- 💻 Multi-class classification (Bacterial vs. Viral).
- 📱 Edge deployment for resource-limited settings.

Impact: Transfer learning with careful fine-tuning can achieve radiologist-level performance, offering scalable diagnostic assistance worldwide.

Thank You

